

The Bridge Experiment: Does an Informal $\leftrightarrow$ Formal Join Move End-Task Lean Performance?

Jack McCarthy

WikiLean · <https://wikilean.jackmccarthy.org>

with experiments executed by Claude Code agents under direction

Final report — July 25, 2026

Preregistration: `docs/research/BRIDGE-EXPERIMENT.md`,
commit `0d36f2664ab2ebb2c7b80b350d3c9bd820414335` (2026-07-16).

Results log: `docs/research/BRIDGE-RESULTS.md`. All data, code, and run transcripts are preserved in the private repository `Deicyde/wikilean-bridge-experiment` (file map in Section 8). Every number below is recomputable from a file in that repository; the file is named where each table appears.

1 Abstract

We test the hypothesis that easy informal $\leftrightarrow$ formal *joining* — a curated dictionary between mathematical concepts and Mathlib declarations (the WikiLean “Brain”), not merely access to both corpora — improves language-model performance on formalization tasks. Five preregistered arms isolate the join: no tools (A), informal search (B), formal search (C), the joined bridge (D), and both corpora unjoined (E). Grading is mechanical throughout; typecheck alone is never success. Three headline results. (1) On 100 contamination-proof tasks built from Mathlib theorems newer than every index, the bridge arm folds 42% success versus 16% for the unjoined-tools control (McNemar 32 vs. 6 discordant, exact $p < 0.0001$), while the no-tool arm collapses from 59.6% to 20% — quantifying ProofNet memorization; the bridge cuts hallucinated-decl citations to 6.8% versus 17.7–26.3% in all other arms. (2) On third-party retrieval gold (MathlibQR-810, MathlibMPR), a Sonnet agent with the union of bridge + formal-search tools plus a measured-behavior manual sets the best rows we know of on both benchmarks (R@10 0.885; group-recall@10 0.557 vs. the specialist SOTA 0.461). (3) The join does not help premise selection (0.272 vs. 0.453 for formal search) — the preregistered concept $\neq$ premise boundary, now quantified as an 18pp API target.

2 Hypothesis and preregistered design

Hypothesis (operationalized). “Necessary” being unprovable, the preregistration commits to three predictions: P1 — an agent with the *joined* dictionary beats an agent given informal and formal search *separately* ($D > E \Rightarrow$ the join, not corpus access, carries the effect); P2 — the bridge solves at lower cost; P3 — the bridge lifts grounding (fewer hallucinated declaration citations), not just compile rate.

The five arms. Identical model, prompt, and budgets per phase; the tool manifest is the only difference (`bench/run_bridge.py`, manifests in `bench/arms/`); see Table 1.

Models per phase. Tier 1 (statement autoformalization): `claude-haiku-4-5` (model id `claude-haiku-4-5-20251001`) in all five arms. Bridge v2 (retrieval and proving): `claude-sonnet-5` in all agent arms, per an explicit pre-execution decision (`docs/research/BRIDGE-V2-BENCHMARKS.md`, 2026-07-25).

Arm	Tools	Isolates
A <code>no_tools</code>	<code>none (-tools "")</code>	floor / memorization
B <code>informal</code>	Wikipedia + nLab search/fetch, no Lean mapping	informal reasoning alone
C <code>formal</code>	loogle + ripgrep + source read over pinned Mathlib	the LeanSearch-class status quo
D <code>wikibrain</code>	the Brain MCP (<code>brain_bridge/search/cell/transfer/neighborhood/snippets/filter</code> + <code>batch decl_exists</code>)	the join
E <code>B+C unjoined</code>	B's tools AND C's tools, no join	the decisive control

Table 1: The five preregistered arms. Identical model, prompt, and budgets; only the tool manifest differs.

Grading rules. Typecheck is never success by itself (TheoremGraph [1]: 22/24 typechecked, 5/24 correct). Tier-1 success = produced a declaration $\wedge$ zero hallucinated citations $\wedge$ typechecks on the pinned toolchain. The preregistered LLM-judge faithfulness leg was quarantined behind a human-calibration gate and ultimately dropped in favor of graders that existed before the arms did (third-party gold labels and the Lean kernel) — see Section 7 for why. Hallucinated-decl rate is graded against a union oracle (doc-gen4 declaration data $\cup$ verified renames; extractor documented in `bench/score_bridge.py`).

Task sets. Tier 1a: ProofNet# — 371 problems (341 eval + 30 prompt-freezing dev, all scored), pinned to PAug/ProofNetSharp @a8da405fbd1e348a87445c2e562c747b7e26dc8f (MIT; [12]); graded on Lean v4.32.0-rc1 / Mathlib a33a5ccd. Tier 1b: 100 fresh tasks built from theorems merged into Mathlib master 2026-07-03→07-16, provably absent from every arm’s index including the Brain’s decl universe (`bench/data/fresh_tasks.stats.json` records the held-out guarantee); graded on Lean v4.33.0-rc1 / Mathlib 9944fe29. A second independent determinacy annotation (agreement 79%, $\kappa \approx 0.20$ — the annotators apply complementary strictness) defines a 74-task both-determinate primary subset.

Budgets. 30 turns stated in the prompt, 10-minute wall-clock enforced, identical across arms; agents had no in-loop typechecker and were instructed to verify cited names through their tools.

Protocol deviations (recorded in the preregistration before any result existed; numbering as in that document):

1. Turn budget prompt-stated, wall-clock-enforced (the CLI lacks `-max-turns`); turns recorded per row.
2. Arm D talks to the production Worker code served locally over the shipped shard bytes (`bench/arms/local_worker.mts`); snapshot pins echoed per response.
3. Grading uses a persistent REPL (Mathlib loaded once) rather than single-shot elaboration; same pins.
4. No in-loop typechecking for agents this campaign — making hallucinated-decl rate a pure grounding measure.
5. Judge grades excluded from all claims pending human calibration (subsequently superseded by the v2 judge-free design).
6. **Skill-leak seal (deviation 6).** The dev smoke showed tool arms B–E calling the CLI’s built-in `Skill/ToolSearch` (32–56 calls/arm), loading e.g. Mathlib-conventions text that arm A structurally could not reach — an asymmetric contamination. `Skill`, `ToolSearch`, `Agent` were added to the disallowed-tools list *before any eval run*; the 120 contaminated B–E dev runs were quarantined unscored to `bench/data/runs_devleak_2026-07-18/` and re-run clean.
7. **Trace telemetry (deviation 7).** Per-call tool traces (truncated input, result head, error flag) were retained from eval arm C onward, observation-only — nothing any agent sees changed. Consequence: eval arms C/D/E and all fresh runs carry traces; eval A/B carry name-counts only.

3 Tier 1: statement autoformalization (claude-haiku-4-5)

3.1 Tier 1a — ProofNet# ($n = 371/\text{arm}$)

Source: docs/research/BRIDGE-RESULTS.md; recomputable from bench/data/bridge_summary.json + bench/data/runs/. Success = produced $\wedge$ no hallucinated citation $\wedge$ typecheck. Results in Table 2.

metric	A none	B wiki	C formal	D wikibrain	E B+C unjoined
success (folded)	59.6%	57.1%	62.3%	64.2% ^a	60.9%
success_proxy (no-halluc)	76.8%	72.8%	72.2%	84.6%	71.2%
typecheck ok	66.6%	63.6%	81.9%	74.7%	83.3%
halluc-decl rate	10.1%	11.0%	10.6% ^a	5.9%	11.3%
runs w/ hallucination	86	101	103	57	107

^a Corrected for display rounding relative to the markdown report and results log, which printed D success 64.1% and C halluc-decl rate 10.7%; the exact recomputed values are $238/371 = 64.15\%$ and $140/1315 = 10.65\%$. No claim is affected.

Table 2: Tier 1a — ProofNet# ($n = 371/\text{arm}$). Bold marks the best column per row.

McNemar D-vs-E (the preregistered test): 63 vs. 51 discordant, exact $p = 0.30$. Direction favors D; underpowered at this effect size. Two readings matter: arm A at 59.6% with zero tools is substantial ProofNet memorization (the motivation for Tier 1b), and the typecheck row *anti-correlates* with grounding — E leads typecheck (83.3%) while trailing on hallucinations, reproducing TheoremGraph’s typecheck-is-not-a-signal finding at $15\times$ their n .

3.2 Tier 1b — the fresh set ($n = 100/\text{arm}$, contamination-proof)

Same files. Memorization stripped, the field reorders (Table 3, Figures 1 and 2):

metric	A none	B wiki	C formal	D wikibrain	E B+C unjoined
success (folded)	20%	22%	25%	42%	16%
halluc-decl rate	21.2%	17.7%	20.9%	6.8%	26.3%
runs w/ hallucination	54	48	49	23	36

Table 3: Tier 1b — the fresh set ($n = 100/\text{arm}$), built from Mathlib theorems newer than every arm’s index.

- **McNemar D-vs-E: 32 vs. 6 discordant, exact $p = 2.4 \times 10^{-5}$.** Also D-vs-C 28 vs. 11, $p = 0.0095$; D-vs-A 29 vs. 7, $p = 0.0003$. The preregistered hypothesis test is decisive on the held-out set: the *join* carries the effect, not tool volume.
- On the 74-task both-determinate primary subset (recomputed from bench/data/fresh_tasks.jsonl det fields $\times$ the paired matrix): D 31/74 (41.9%) vs. E 14/74 (18.9%), McNemar 22 vs. 5, exact $p = 0.0015$ — the effect is not an artifact of underdetermined prompts.
- Arm A collapses $59.6\% \rightarrow 20\%$ off-distribution (memorization quantified).
- D’s hallucination advantage *widens* off-distribution (6.8% vs. 17.7–26.3%).
- Arm E is worst at 16% and produced no declaration at all in 31/100 fresh runs — unjoined tool volume can be actively harmful.
- Cost (recomputed from run rows; mean per task, folded over all 471 runs/arm): A \$0.034 · B \$0.048 · C \$0.140 · **D \$0.121** · E \$0.128; mean wall-clock C 116 s · D 89 s · E 106 s. The bridge arm outscores the cheaper-per-task C and E while costing less than either — P2 holds directionally.

3.3 Trace attribution (deviation-7 telemetry, eval C/D/E)

From bench/trace_analysis.py over bench/data/runs/ (eval split):




Figure 1: Tier-1 folded success by arm, ProofNet# eval versus the contamination-proof fresh set. Off-distribution, the no-tool arm collapses 59.6% $\rightarrow$ 20% while the bridge arm (D) holds 42% against 16% for the unjoined-tools control (E). Data: `bench/data/bridge_summary.json` (paired matrix, split by task id).




Figure 2: Hallucinated-declaration citation rate by arm (share of cited declaration names absent from the union oracle). D’s grounding advantage *widens* off-distribution: 6.8% on the fresh set versus 17.7–26.3% for every other arm. Data: recount over `bench/data/runs/` with the `bench/score_bridge.py` extractor and oracle.

- 98–100% of traced tool-arm runs cite at least one declaration that visibly surfaced in a tool result; only 10–17% of citations never touched the tools.
- **35% of arm-D runs (116/329) cite a declaration that came out of a brain_bridge / brain_transfer result** — English in, formal name out. This is an undercount: result heads truncate at 200 chars.
- Among *hallucinated* citations, the checked-and-cited-anyway rate is C 35%, E 30%, **D 13%** — binary `decl_exists` verification disciplines the model where fuzzy search neighborhoods (loogle/grep) let it fool itself. This is the mechanism behind the halluc rows above.

4 Third-party retrieval: MathlibQR-810 and MathlibMPR (claude-sonnet-5)

Motivation: Tier-1’s remaining faithfulness leg depended on an LLM judge whose error could correlate with arm. Bridge v2 replaces judge-dependent grading with graders that predate the arms: third-party expert gold labels here, the Lean kernel in Section 5. Data: `frenzymath/LeanSearch-v2 @94f4888cbaf9`, CC BY 4.0 (LeanSearch-v2 [3]), copied under `bench/v2/data/`. Arms: N (no tools), F (loogle + `decl_grep` + `decl_read`), W (the Brain MCP), WF ($W \cup F$ **plus** `bench/v2/AGENT_MANUAL.md` prepended to the prompt — the manual is deliberately part of the condition; see the confound note below). Scoring: `bench/v2/score_retrieval.py` (exact full-name match only); run rows with full gzipped stream-json transcripts in `bench/v2/runs/agent/`.

4.1 MathlibQR fair-810 — concept retrieval (810 expert query rows, 6 styles)

system	R@10	nDCG@10
published: TheoremGraph	0.775	0.548
published: LSv2 retriever+reranker	0.780	0.623
system-mode wikibrain (one <code>brain_bridge</code> call, no LLM)	0.036	0.031
agent N	0.633	0.598
agent F	0.831	0.790
agent W	0.816	0.781
agent WF	0.885	0.839

Table 4: MathlibQR fair-810 concept retrieval. Exact full-name match; gold is one declaration per query row.

- **The label-resolver finding.** System-mode scores 0.036 — the API’s free-text entry is a label/alias resolver, not a semantic retriever. Nickname queries reach nDCG 0.196 while the descriptive styles (LaTeX, natural-language, special-case) score 0.000. Yet an agent *iterating* over the same API reaches 0.816 — the content is in the graph; the single-shot entry point is the gap. This is the headline API deficiency.
- **F vs. W: statistically tied** (paired discordants 78 F-only vs. 66 W-only, exact $p = 0.36$); both nominally above the published retriever rows, with the apples-to-oranges caveat that agents reason and verify while retrievers embed once (Section 7).
- **Style texture:** W wins the `special_case` style (nDCG 0.523 vs. F’s 0.384 — the Brain’s curated `special_case` bonds are the signal); F wins Lean-syntax styles. Grep cannot find what source text never names.
- **WF is decisive:** beats F (71 vs. 27 discordant, exact $p = 1.0 \times 10^{-5}$) and W (83 vs. 27, $p = 8.3 \times 10^{-8}$); +10.5pp R@10 over the best published retriever row. The style spread narrows: every style ≥ 0.55 nDCG, five of six ≥ 0.81 .
- Agent N’s 0.633 includes 143/810 rows scored 0 for format non-compliance (no JSON array in the final message) — an agent-mode failure mode the manual explicitly addresses.




Figure 3: Retrieval by system, with published anchors as dashed reference lines. Left: MathlibQR fair-810 R@10 (anchors: TheoremGraph 0.775, LeanSearch-v2 retriever+reranker 0.780). Right: MathlibMPR group-recall@10 (anchors: LeanSearch-v2 reasoning 0.461, DIVER 0.380). System-mode is a single `brain_bridge` call with no LLM; N/W/F/WF are Sonnet agent arms. Data: rerun of `bench/v2/score_retrieval.py`.

- Tool census: F mean 2.2 calls/run; W 3.5 (`decl_exists` 1,251 + `brain_bridge` 608 — verify-then-cite); WF mean 4.6 calls pooled across both benchmarks, a genuine dual-toolkit mix (`decl_grep` ~1.5k, `decl_exists` ~0.9k, `brain_bridge` ~0.7k, `decl_read` ~0.5k, `loogle` ~0.4k). The `brain_cell` misuse pattern (63% failure rate in Tier-1 traces) disappears under the manual.

4.2 MathlibMPR — premise retrieval (69 post-cutoff PR theorems, expert premise groups)

system	group-recall@10
published: LeanSearch-v2 reasoning	0.461
published: DIVER	0.380
published: TheoremGraph concept-tuned	0.165
system-mode wikibrain	0.000
agent N	0.203
agent W	0.272
agent F	0.453
agent WF	0.557

Table 5: MathlibMPR premise retrieval: a group is covered when any of its interchangeable members appears in the top 10.

- A generic Sonnet agent with `grep` + `loogle` **matches the specialist premise retriever** (0.453 vs. 0.461) — a notable baseline result on its own.
- The Brain helps over memory (+7pp) but trails formal search by 18pp: the preregistered concept $\neq$ premise boundary (TheoremGraph’s negative transfer), now measured on our own tools. `brain_premises` (BRIDGE-ISSUES #7) has a quantified 18pp target.
- WF beats both the best single arm (+10pp over F) and the published SOTA on its own benchmark (0.557 vs. 0.461).

The WF confound, stated plainly. WF = union tools + the evidence-based manual. The two are not separated in this campaign; a bare-union ablation is one flag away (`MANUAL_ARMS` in `bench/v2/run_agent.py`) and was not run. WF results should be read as “a maximally equipped, briefed agent”, not as a clean marginal effect of tool union.

5 SorryDB: kernel-graded proving in live repositories (claude-sonnet-5)

SorryDB [5] serves unsolved `sorrys` from real Lean repositories; success is defined by the repo *building* with your proof. No judge exists anywhere in the loop and no solutions exist to memorize. Snapshot: the SorryDB-2601 1,000-task evaluation split (`bench/v2/data/sorrydb/SorryDB_2601_1000_evaluation_split.json`).

Subset construction and the snapshot-rot finding. Disk permitted one repository build at a time, so we sampled repositories rather than tasks: every task of a chosen repository is included, and verification cost amortizes across it. The original freeze — 8 repositories, 164 tasks, chosen for toolchain availability, build weight, and domain diversity, plus one pedagogical floor — lost 74 of 164 tasks (45%) in the roughly six months since the snapshot: the pinned commits of GlimpseOfLean, LeanCourse25, and most Foundation and SemicircleLaw tasks no longer exist on GitHub. When a repository’s history is rewritten, commits no branch reaches are eventually deleted, and a snapshot’s pins die with them. We could not retrieve these commits by any route (direct git fetch by hash, GitHub’s archive downloads, or the authenticated API). Detection is the subtle part: the archive endpoint answers preflight HEAD requests for dead commits with HTTP 200, so a liveness check that does not attempt a real fetch reports the pins as healthy. We record this as a methodological finding: benchmarks pinned to live repositories rot at the commit level, and the rot is invisible unless every pin is actually re-fetched before use. The refreeze (`bench/v2/sorrydb_prep.py`, commit `db679e8d`) keeps only verified-fetchable pins: 171 tasks across 10 repositories (PrimeNumberTheoremAnd 21, curve25519-dalek-lean-verify 21, category-theory-in-context 21, PersistentDecomp 20, brownian-motion 20, graphiti 20, LeanAPAP 20, MiscYD 20, SemicircleLaw 7, Foundation 1).

Protocol. Agent phase is build-free and one-shot: the row carries the goal state plus a ± 60 -line file window at the pinned commit; the agent (600s timeout, arms N / F / WF as in Section 4) outputs one replacement for the `sorry`, with an explicit honest-abstention rule (output literal `sorry`). Verification (`bench/v2/verify_sorrydb.py`) clones each repo at its pin, establishes a pristine baseline build, splices each candidate over the span (asserted to literally read `sorry`), and rebuilds the module — success iff exit 0 with no `sorry/axiom` warning.

Results — recomputed for this report from `bench/v2/runs/sorrydb/verify.jsonl` + the run rows (script logic in Section 8; $n = 171$ tasks/arm) — in Table 6 and Figure 4.

Notes, in the interest of a complete record: (i) 8 candidate proofs — all from curve25519-dalek-lean-verify, whose verification pass covered only 2 of its 10 candidates — carry no kernel verdict in `verify.jsonl`; they are conservatively counted as not-proved above. (ii) 3 (N) and 2 (F) tasks have no run row at all (runner losses). (iii) 2 stale verdicts against dead-pin rows from the first verification pass are excluded. Distinct sorries proved across all arms: 13; N’s proofs are a strict subset of both F’s and WF’s; WF vs. F overlap 6, WF-only 4, F-only 3 — no significance claim at this n .

Reading. Tools triple-to-quintuple the kernel-verified prove rate and nearly triple the honest-abstention rate (N instead burns its budget: 70/168 rows produced nothing, mostly 600s timeouts). Cost-per-proved-theorem *falls* by $2.7\times$ as tool cost rises — verification effort is cheaper than blind search. **Comparison caveat:** the published SorryDB agentic best ($\approx 30\%$, union 35.7%) is not comparable — different task subset (ours is weighted toward research-frontier repos: PNT+, a cryptographic verification codebase, LeanAPAP) and a materially different protocol (published agents iterate against the build; ours drafted one-shot with no compiler in the loop).

	N no tools	F formal	WF union+manual
run rows present	168	169	171
no output (timeout / no lean block)	70	15	11
honest give-up (literal <code>sorry</code>)	40	83	86
candidate proofs submitted	58	71	74
proved (kernel)	2	9	10
failed	48	54	56
unsplicable	5	6	5
no verdict (see note)	3	2	3
proved / 171 tasks	1.2%	5.3%	5.8%
proved / candidates	3.4%	12.7%	13.5%
total cost (USD)	\$58.97	\$102.76	\$108.87
cost per proved	\$29.48	\$11.42	\$10.89
mean wall-clock / task	120 s	198 s	183 s

Table 6: SorryDB kernel-graded proving, one-shot with no compiler in the loop ($n = 171$ live-pin tasks/arm).




Figure 4: SorryDB, one-shot drafting with no compiler in the loop. Left: kernel-verified proofs out of 171 live-pin tasks. Right: cost per proved theorem — tools raise total spend but *cut* the cost of each success by $2.7\times$. Data: `bench/v2/runs/sorrydb/verify.jsonl` $\times$ the run rows.

6 Synthesis

What the join adds. Two things, both now measured twice. *Grounding*: binary existence verification at the informal→formal boundary collapses hallucinated citations (5.9%/6.8% vs. 10–26%; checked-and-cited-anyway 13% vs. 30–35%) — and the effect grows off-distribution, where memory fails. *Concept entry*: when the query is a description rather than a name — the fresh set, the special_case query style — the curated join finds what lexical and type-pattern search cannot (42% vs. 16–25%; nDCG 0.523 vs. 0.384). The fresh-set McNemar ($p < 0.0001$) says the join itself, not corpus access, carries this.

Where it does not. Premise selection ranks by a different signal than concept identity: W trails F by 18pp on MathlibMPR, exactly as preregistered from TheoremGraph’s negative transfer. And the API’s single-shot free-text entry is a label resolver (0.036 R@10 system-mode) — everything the agent arms recover, they recover by reformulating against it.

The composition result. The toolkits are complementary, and a briefed agent composes them into the best rows we know of on both third-party benchmarks (QR 0.885, MPR 0.557) — on gold labels we did not write. The Bridge thesis restated in retrieval form: the join *plus* the territory beats either alone.

API roadmap, derived. (1) A semantic statement-level index behind `brain_bridge` — the 0.036 → 0.816 single-shot/agent gap is an entry-point problem, not a content problem. (2) `brain_premises(goal_state)` — a premise-level ranking mode with a measured 18pp target (BRIDGE-ISSUES #7). (3) The eight pre-campaign API requirements (batch `decl_exists`, signatures + imports per hit, generalization surfacing, honest abstention, censored neighborhoods, snapshot echo, the composite `brain_bridge`, teaching tool descriptions) were implemented before the campaign and are what the arms exercised. (4) The `AGENT_MANUAL.md` pattern itself: measured tool behavior (~5,500 logged calls) distilled into a briefing eliminated the dominant misuse mode (`brain_cell` 63% → ~0) and is the cheapest intervention in the whole study.

Cost economics. Tier-1: the bridge arm is cheaper *and* better than its tool-bearing competitors (\$0.121/task vs. C \$0.140, E \$0.128, at +2–26pp success). Retrieval: \$0.08–0.21/query (QR), \$0.18–0.44/query (MPR). Proving: tool arms cut cost-per-proved-theorem from \$29 to \$11 while raising the prove rate 5×. Across all three phases, adding the join never made the task more expensive per success.

7 Limitations and threats to validity

1. **One model per phase.** Tier-1 is Haiku-only (per-model run dirs were not keyed — fixed in v2, where dirs are (benchmark, arm, model)-keyed); v2 is Sonnet-only. A bridge that only helps some capability band would not be distinguished. The preregistered two-model grid remains undone.
2. **Agent-vs-retriever comparability.** Our QR/MPR agent rows spend multi-call reasoning and verification per query; published retriever rows embed once. “Above the published rows” is a statement about achievable quality at agent cost, not a same-budget comparison. System-mode is the same-budget comparison, and it loses badly.
3. **The WF manual confound.** WF bundles tool union with the briefing; the ablation exists in the harness but was not run.
4. **Power on secondary contrasts.** F-vs-W on QR is a null at $n = 810$ ($p = 0.36$); MPR has $n = 69$; SorryDB proved counts are single digits per arm — overlap patterns there are descriptive only.
5. **The judge leg was dropped, not repaired.** Preregistered faithful@budget (BEq+ equivalence) was never graded: the judge required human calibration (deviation 5), Jack’s arm-correlated-bias critique stood, and the field precedent (TheoremGraph rejecting its own judge as over-generous) argued against resting headline claims on one. The v2 replacement — graders that predate the arms — is stronger, but it means Tier-1 “success” can still reward a well-formed wrong statement; the Tier-1 fresh-set numbers are best read jointly with the hallucination and trace evidence.

6. **Tier-1a is contaminated by construction** (arm A: 59.6%); it is reported for the paired deltas and as the memorization control against Tier-1b, never alone.
7. **SorryDB caveats.** One-shot drafting without a compiler under-represents every arm relative to published loops; the subset is repo-sampled, not difficulty-matched; 8 candidates lack verdicts; and 45% snapshot rot means the *original* frozen subset is already partly unreproducible upstream — the preserved `tasks_frozen.jsonl` (goal states + context windows) is the durable record.
8. **Mechanical extractors.** Cited-declaration extraction is a documented regex heuristic; hallucination rates are comparable across arms (same extractor) but not exact.

8 Data availability

Everything needed to recompute this report lives in the private preservation repository **Deicyde/wikilean-bridge-experiment** (report at the root as `README.md` and under `report/`). Provenance commits refer to the WikiLean working repository. See Table 7.

Path (preservation repo)	Contents	Key provenance commits (WikiLean)
<code>docs/research/BRIDGE-EXPERIMENT.md</code>	preregistration incl. deviations 1–7	0d36f266 (2026-07-16)
<code>docs/research/BRIDGE-RESULTS.md</code>	Tier-1 + retrieval results log	53f5cb31, daac6107
<code>docs/research/BRIDGE-V2-BENCHMARKS.md</code>	v2 design + verified sources	—
<code>docs/research/BRIDGE-ISSUES.md</code>	issue queue / history	cdc2d743
<code>bench/*.py,</code> <code>bench/README.md,</code> <code>bench/arms/</code>	Tier-1 harness: runner, scorer, REPL typecheck rig, trace analysis, arm manifests	068abe34, e554902e, 6c2ca704, 31b95caa, 79ac3dfe
<code>bench/data/bridge_tasks.jsonl</code> (+.stats)	371 ProofNet# tasks; source pin + MIT licence in stats	e554902e, ec2a9e11
<code>bench/data/fresh_tasks.jsonl</code> (+.stats)	100 fresh tasks + determinacy annotations + held-out guarantee	df5dbf92, a0d45103
<code>bench/data/gold_census.json,</code> <code>fresh_census.json</code>	which golds elaborate on which pin; fresh pin = Lean v4.33.0-rc1 / Mathlib 9944fe29 (Tier-1a pin = v4.32.0-rc1 / a33a5ccd, recorded in <code>bench/README.md</code>)	01e8612b, 185377cc
<code>bench/data/bridge_summary.json</code>	Tier-1 scored summary: per-arm aggregates, paired matrix, McNemar tables	53f5cb31
<code>bench/data/runs/</code> <code>bench/data/runs_devleak_2026-07-18/</code>	all 2,355 Tier-1 run rows (5 arms × 471) the 120 quarantined skill-leak dev runs (deviation 6 evidence)	— f89e7a41
<code>bench/v2/</code> <code>bench/v2/data/</code>	v2 harness + <code>AGENT_MANUAL.md</code> MathlibQR/MPR (LeanSearch-v2 @94f4888cbaf9, CC BY 4.0), SorryDB-2601 split, <code>tasks_frozen.jsonl</code>	c7629584, 21032c06 c7629584, 87638cfc, db679e8d
<code>bench/v2/runs/</code>	every v2 run row with full gzipped stream-json transcripts (QR 810×4, MPR 69×4, SorryDB 508 rows) + <code>verify.jsonl</code> (197 kernel verdicts)	3664aa3d, daac6107, 382e51bf, 1fe21cca

Table 7: File map of the preservation repository.

Recomputation entry points: `bench/score_bridge.py` (Tier-1), `bench/trace_analysis.py` (Section 3.3), `bench/v2/score_retrieval.py` (Section 4 — reprints every table), `bench/v2/verify_sorrydb.py` (Section 5 verdicts; the Table 6 rows are a straight aggregation of `verify.jsonl` × the run rows, filtered to `tasks_frozen.jsonl` ids, counting `gave_up`, `error`, verdicts, and `transcript_stats.cost_usd`). The figures in this document are generated from those same files by `report/bridge-report/recompute_figdata.py` (which asserts every plotted value against the markdown report) and `generate_figures.py`.

External sources cited (all fetched and verified during the v2 design sweep): TheoremGraph [1] · LeanSearch [2] · LeanSearch-v2 [3] · LeanExplore [4] · SorryDB [5] · FATE [6] · miniCTX [7] · LeanDojo [8] · LeanAgent [9] · Numina-Lean-Agent [10] · miniF2F-v2 [11] · ProofNet#/ProofNetVerif [12] · benchmark-faults survey [13].

References

- [1] TheoremGraph. arXiv:2606.25363. <https://arxiv.org/abs/2606.25363>
- [2] LeanSearch. arXiv:2403.13310. <https://arxiv.org/abs/2403.13310>
- [3] LeanSearch-v2. arXiv:2605.13137. Data: pinned frenzymath/LeanSearch-v2 @94f4888cbaf9, CC BY 4.0. <https://arxiv.org/abs/2605.13137>; <https://github.com/frenzymath/LeanSearch-v2>
- [4] LeanExplore. arXiv:2506.11085. <https://arxiv.org/abs/2506.11085>
- [5] SorryDB. arXiv:2603.02668. <https://arxiv.org/abs/2603.02668>; <https://sorrydb.org>
- [6] FATE. arXiv:2511.02872. <https://arxiv.org/abs/2511.02872>
- [7] miniCTX. arXiv:2408.03350. <https://arxiv.org/abs/2408.03350>
- [8] LeanDojo. arXiv:2306.15626. <https://arxiv.org/abs/2306.15626>
- [9] LeanAgent. arXiv:2410.06209. <https://arxiv.org/abs/2410.06209>
- [10] Numina-Lean-Agent. arXiv:2601.14027. <https://arxiv.org/abs/2601.14027>
- [11] miniF2F-v2. arXiv:2511.03108. <https://arxiv.org/abs/2511.03108>
- [12] ProofNet#/ProofNetVerif. arXiv:2406.07222. Tasks pinned to PAug/ProofNetSharp @a8da405f (MIT). <https://arxiv.org/abs/2406.07222>
- [13] Benchmark-faults survey. arXiv:2606.29493. <https://arxiv.org/abs/2606.29493>